

31251 – Data Structures and Algorithms

Week 9, Autumn 2020 - Graphs Part II

Xianzhi Wang

- Dynamic Programming
 - Some Examples
 - Dijkstra's Algorithm
- Connectivity in Graphs
 - Connectivity and Connected Components
 - Strongly Connected Components
 - Articulation Points
- Dependencies in Directed Graphs

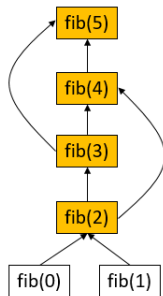
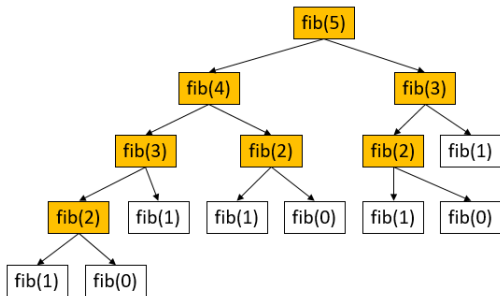
Dynamic Programming

- **Dynamic programming** looks a lot like Divide-and-Conquer, but subinstances overlap.
- Usually implemented with some sort of table that keeps track of subinstance solutions.
- Normally they work bottom-up (Can be done top-down recursively):
 - Solve the smallest subinstances, combine them to make bigger subinstances, keep going until you have the whole thing.

- **Optimal Substructure:** same as greedy algorithms and D&C.
 - $Solution(P) \leftarrow f(Solution(SP_1), \dots, Solution(SP_n))$
- **Overlapping Subproblems:**
 - e.g., the recursive algorithm (practice 4a, week 4) repetitively calculates the subinstances of the problem.

```
fib(5)
= fib(4) + fib(3)
= (fib(3) + fib(2)) + (fib(2) + fib(1))
= ((fib(2) + fib(1)) + (fib(1) + fib(0))) + ((fib(1) + fib(0)) + fib(1))
= (((fib(1) + fib(0)) + fib(1)) + (fib(1) + fib(0))) + ((fib(1) + fib(0)) + fib(1))
```

Recursive (on the left) vs. Dynamic Programming (on the right)



“Given n items $\{o_1, o_2, \dots, o_n\}$ and their weights $\{w_1, w_2, \dots, w_n\}$ and values $\{v_1, v_2, \dots, v_n\}$, put these items in a knapsack of capacity W to get the maximum total value in the knapsack.”¹

- Overlapping subproblems:
 - Consider all possible solutions, especially those near-to-optimal ones, they are likely to contain similar subsets of items.

- Optimal substructures:

- If the optimal solution contains only one item:

$$optsol_1 = \arg \max_{i \in \{1, 2, \dots, n\}} \{v_i | w_i < W\}$$

- If the optimal solution contains m items:

$$optsol_m = \arg \max_{i \in \{1, 2, \dots, n\} / sol_{m-1}} \{v_i + v(sol_{m-1}) | w_i + w(sol_{m-1}) < W\}$$

¹<https://www.geeksforgeeks.org/0-1-knapsack-problem-dp-10/>

Given G with non-negative edge weights and a starting vertex u :

- ① Set distances: 0 for u , $+\infty$ for others vertices.
- ② Mark all vertices as unvisited.
- ③ Set the current vertex to u .
- ④ While there are unvisited vertices:
 - ① For each unvisited neighbour of the current vertex:
 - ① Compare the distance of each neighbour to the distance to current plus the edge weight of the edge joining them.
 - ② Keep whichever is smaller.
 - ② Mark the current vertex as visited.
 - ③ Select the unvisited vertex with smallest tentative distance and set it as the current vertex.

- Where was the dynamic programming?

Suppose $dist[v]$ is the distance from u to v , then the algorithm is repeatedly performing

$$\underbrace{dist[v]}_{\text{updated distance}} = \min \left\{ \underbrace{\overbrace{dist[v]}^{\text{past result}}}_{\text{known shortest distance}}, \underbrace{\overbrace{dist[current]}^{\text{past result}} + w(current, v)}_{\text{newly calculated distance}} \right\}.$$

- As such, we always keep the shortest distance of a vertex from the starting vertex, representing an optimal substructure.

- Suppose $|V| = n$ and $|E| = m$ in G . We have $n - 1 < m < n(n - 1)$.
- There are $O(n)$ inserts, $O(m)$ updates, $O(n)$ Find/Delete Minimums
- Data structure-dependent:
 - Unsorted array: $O(n^2)$.
 - Binary min heap: $O(m \log n)$.
 - Fibonacci heap²: $O(m + n \log n)$

²https://en.wikipedia.org/wiki/Fibonacci_heap

Dealing with Negative Edges

- Dijkstra's algorithm ran into trouble with negative edges because we never look at marked vertices again.
- Bellman-Ford algorithm can handle negative edges – it just looks at everything again!
- Doesn't work with negative *cycles* though – but it can detect them.

Bellman-Ford Algorithm:

- ① Set things up as with Dijkstra's.
- ② For 1 to the number of vertices
 - For each edge uv
 - If using uv improves the distances, use it and update as needed.

Check for negative cycles:

- ③ For each edge uv
 - If uv improves the distances, there's a negative cycle.

Suppose $|V| = n$ and $|E| = m$ in G .

Bellman-Ford gets:

- $O(nm)$ time – can't really get around this.
- $O(n)$ space.

Connectivity

- When using graphs to model things, we usually want to know if they're all in one piece.
- A **connected** graph allows any vertex to get to any other vertex via edges.
- It is useful to check the reachability between nodes in computer networks or transport networks, etc.

- Testing connectivity is easy — Just do a traversal!
 - If we visit all the vertices, the graph must be connected.
- What if it's not connected?
 - We might be interested in finding out what the subgraphs are.
 - We call them **connected components**.
- Finding them is almost as easy as testing connectivity:
 - ① While not all vertices have been visited:
 - ① Pick an unvisited vertex to start a new traversal.
 - ② Mark all the vertices you can reach as visited and record them as a component.

Strongly Connected Components

For directed graphs,

- Two vertices are *strongly connected* if there is a path from one to the other, as well as a path back.
- *Strongly connected component*: a subgraph where every pair of vertices is strongly connected.

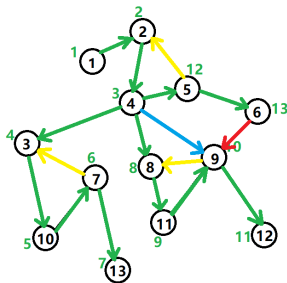
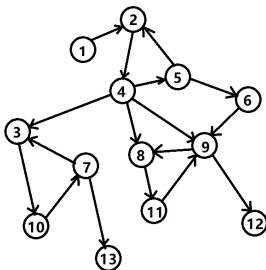
Note that,

- In a directed graph, a may not be adjacent to b when b is adjacent to a .

Detecting Strongly connected components in a graph:

- We can just run a bunch of reachability queries.
 - $O(n^2(n + m))$ without clever optimisation.
 - $O((n + m) \log n)$ if we make the approach complicated enough (uses divide and conquer).
- Parallelises well.

We can do it faster using *Tarjan's Algorithm*.



DFS spanning trees have four kinds of edges³:

- Tree edge (green): *points to* an unvisited node.
- Back edge (yellow): *points to* an ancestor.
- Cross edge (red): *points to* a visited non-ancestor.
- Forward edge (blue): *points to* an visited offspring.

³<https://programmer.ink/think/5d2cd3c392a88.html>

- Each vertex has an $\text{id} \in \{1, 2, \dots, n\}$.
- Each vertex u maintains two variables:
 - $\text{DFN}[u]$: The order in which nodes are searched during depth-first search traversal.
 - $\text{LOW}[u]$: Let the subtree rooted in u be $\text{Subtree}(u)$. $\text{LOW}[u]$ is defined as the minimum value of the following nodes: the node in $\text{Subtree}(u)$; the node from $\text{Subtree}(u)$ through an edge not on the search tree.
- An ancestor has smaller DFN than offspring in the tree.
- DFN on a path starting from the root increases strictly while LOW does not.

- Nodes in the graph are searched and SCCs are detected during depth-first-search (DFS) traversal.
- Given the current vertex u , consider three cases its adjacent node v (not the parent of u):
 - **Case 1: v is not accessed**: continue to search for v in depth. In the retrospective process, update $LOW[u]$ with $LOW[v]$ — whatever v can trace back to, u can as well.
 - **Case 2: v has been visited and is in the stack**: v could be the earliest vertex that can be traced back from u . Update $LOW[u]$ with $DFN[v]$.
 - **Case 3: v has been visited but is not in the stack**: v has been added to a connected component. There is nothing left to do.

Tarjan's Algorithm

```
int dfn[n] = {0}, low[n], dfn_cnt = 0; // variables for each vertex
stack<int> s; // stack
set<int> in_s; // a set for checking if a node is in stack
set<set<int>> all_sccs; // algorithm's result: all SCCs

void tarjan(int u) {
    low[u] = dfn[u] = ++dfncnt; // initialise dfn and low
    s.push(u); in_s.insert(u); // add u to stack & mark it as visited
    for (v: every adjacent vertex of u) {
        if (!dfn[v]) // case1: v unvisited
            tarjan(v); // recur for v
            low[u] = min(low[u], low[v]); // retrospective
        } else if (in_s.find(v) != in_s.end()) // case2: v visited & in stack
            low[u] = min(low[u], dfn[v]); // update low[u]
    }
    if (dfn[u] == low[u]) { // this only happens when u is the root of a SCC
        set<int> scc; // create a new scc
        while (s.top() != u) // add all vertices above u in stack to scc
            scc.insert(s.top()); in_s.erase (s.top()); s.pop();
        scc.insert(s.top()); in_s.erase (s.top()); s.pop(); // add u
        all_sccs.insert[scc]; // put scc to the result
    }
}
```

```
int dfn[n] = {0}, low[n], dfn_cnt = 0; // variables for each vertex
stack<int> s; // stack
set<int> in_s; // a set for checking if a node is in stack
set<set<int>> all_sccs; // algorithm's result: all SCCs

void find_scc(directed_graph d) {
    for (each vertex u in d) {
        if (!dfn[u]) // u unvisited
            tarjan(u);
    }
}
```

Complexity:

- Time: $O(n + m)$.
- Space: $O(n)$

A similar algorithm can be used to find the equivalence of strongly connected components in an undirected graph.

- “A vertex in an undirected connected graph is an **articulation point** (or **cut vertex**) *iff* removing it (and edges through it) disconnects the graph.”⁴
- “Articulation points represent vulnerabilities in a connected network — single points whose failure would split the network into two or more disconnected components.”

⁴<https://www.geeksforgeeks.org/articulation-points-or-cut-vertices-in-a-graph/>

- Pick an arbitrary vertex of the graph root and run depth first search from it.
 - Undirected graphs have only tree and back edges (why?)
 - All vertices visited before u in DFS are ancestors of u .
- Given the current vertex in DFS, v ($\neq \text{root}$),
 - Case 1: if none of its descendants has a back-edge to its ancestors, v is an articulation point;
 - otherwise, v is not an articulation point.
- Case 2 (special case): *root* is an articulation point *iff* it has more than one child in the DFS tree.
 - Why? That is at least one subtree that can only be reached (by other vertices) through *root*.

⁵<https://cp-algorithms.com/graph/cutpoints.html>

- Each vertex has an id $\in \{1, 2, \dots, n\}$.
- Each node u maintains two variables:
 - $tin[u]$: the entry time (*or* age) for u .
 - $low[u]$: the lowest entry time reachable by u ⁶.

$$low[u] = \min \begin{cases} tin[u] & \text{the entry time of } u \\ tin[v] & \text{for all } v \text{ for which } (u,v) \text{ is a back edge} \\ low[to] & \text{for all } to \text{ for which } (u,to) \text{ is a tree edge} \end{cases}$$

- u in the DFS tree is an articulation point *iff* its depth is smaller than or equal to the lowpoint of any of its children
 - $tin[u] \leq low[to]$

⁶ $low[]$ is just a way of tracking whether there are at least two paths to get to a vertex — Vertices in biconnected components all have the same low value.

```
int tin[n] = {0}, low[n], timer = 0;
bool visisted[n] = {false};

void find_cutpoints(undirected graph d) {
    for (u: every vertex in d) {
        if (!visited[u])
            dfs(u);
    }
}
```

```
int tin[n] = {0}, low[n], timer = 0;
bool visisted[n] = {false};
set<int> cutpoints;

void dfs(int u, int v = -1) { // no parent of u is known by default
    visited[u] = true;
    low[u] = tin[u] = ++timer;
    int childCount = 0;
    for (int to : all children of u) {
        if (to == v) // (u,to) leads back to u's parent in DFS tree
            continue;
        if (visited[to]) { // (u,to) is a back edge to an ancestor
            low[u] = min(low[u], tin[to]);
        } else { // u unvisited, (u,to) is a tree edge
            dfs(to, u);
            low[u] = min(low[u], low[to]);
            if (low[to] >= tin[u] && v != -1) // case 1
                cutpoints.insert(u);
            ++childCount;
        }
    }
    if (v == -1 && childCount > 1) // case 2
        cutpoints.insert(u);
}
```

```
FindArticulationPoints(i, d)
  visited[i] = true, depth[i] = d
  low[i] = d, child_count = 0
  is_articulation = false
  for (each neighbour n of i)
    if (not visited[n])
      parent[n] = i
      FindArticulationPoints(n, d+1)
      child_count = child_count + 1
      if (low[n] >= depth[i])
        is_articulation = true
        low[i] = min{low[i], low[n]}
      else if (n != parent[i])
        low[i] = min{low[i], depth[n]}
  if ((i has a parent and is_articulation) or
      (i has no parent and child_count > 1))
    i is an articulation point
```

- Time: $O(n + m)$ — it's just a depth-first traversal.
- Space: $O(n)$
- Can be modified to retrieve the biconnected components with the same complexity.
- What if we just remove each edge and test if the remaining graph is still connected?
 - $O(n * (n + m))$
- Can be done in parallel in $O(\log n)$ time! (But with $O(n + m)$ processors...)

Dependencies in Directed Graphs

If a directed graph has no directed cycles, we called it a *Directed Acyclic Graph*, or *DAG*.

How many strongly connected components does this have?

DAGs are like the trees of the directed graph world.

As such, they have many applications:

- Anything where dependencies are important (anyone running a *NIX distro will understand this one).
- Task scheduling.
- Compilation.
- Computing information in a spreadsheet.
- Modelling multi-stage processes.

DAG describes the dependencies among things (vertices).

We are interested in how to do things in an order that is compliant with the dependencies — *topological ordering*.

It may not be unique. (why?)

We do *topological sorting* to find a topological ordering.

```
list<vertex> topo_order;

void Kahn_topo_sort(directed graph g) {
    set<vertex> s = {all vertices with no incoming edges};
    while (!s.empty()) {
        v = s.remove_something_from_a_set();
        topo_order.append(v);
        for (each neighbour u of v) {
            remove (v,u) from g;
            if (u has no incoming edges)
                add u to s;
        }
    }
}
```

Depth-First Topological Sort

```
stack<vertex> topo_order;

void topo_sort (directed graph g) {
    while (not all vertices marked) {
        v = next unmarked vertex;
        visit(v)
    }
}

visit (vertex v) {
    if (v is marked) // terminate the iteration
        return;
    for (each (v,u) in edges) // do recursion
        visit(u);
    mark v;
    topo_order.push(v); // update the result
}
```

- Both linear time algorithms – $O(n + m)$.
- Both use $O(n)$ extra space.
 - Can get $O(1)$ for Kahn's algorithm at the cost of increasing time to $O((n + m) \log n)$.
- Can find a topological (a different way) in $O(\log^2 n)$ time ... if we have a polynomial number of processors.

Appendices

Dijkstra's Algorithm – Correctness

Let S be the set of marked vertices, u the starting vertex, and $d(v)$ the calculated distance from u to v , then the correctness of Dijkstra's algorithm can be stated as:

Lemma

For every vertex $v \in S$ $d(v)$ is the shortest path from u to v .

Dijkstra's Algorithm – Correctness

Proof. By induction on $|S|$.

- **Base Case.** Easy, $|S| = 1$ implies $S = \{u\}$, and the distance is 0. Alternatively $|S| = 0$ is trivial.
- **Induction Hypothesis.** Assume the theorem is true for $|S| = k \geq 1$.
- **Inductive Step.**
 - Let v be the next vertex added to S , and w be the neighbouring path vertex according to the algorithm.
 - By the assumption, $d(w)$ is minimum.
 - Let P , for contradiction, be the shortest path from u to v .
 - Let xy be the first edge in P that leaves S .
 - Then the length of P is at least $d(x) + w(x, y)$, which is the tentative distance to y .
 - This distance can't be smaller than the tentative distance to v , otherwise we would've chosen it instead of v to add next.
 - Therefore the length of P is at least $d(v)$. □

Formally, we need a slightly more precise statement:

Lemma

After i iterations of the outer for loop:

- *If $d(v) \neq \infty$, there is a (known) path from the start vertex u to vertex v with distance $d(v)$.*
- *If there is a path from u to v with at most i edges, $d(v)$ is the length of the shortest path with at most i edges.*

Bellman-Ford Algorithm - Correctness

Proof. By induction on i :

- **Base Case:** $i = 0$, nothing has happened yet, trivial!
- **Inductive Assumption:** Assume that at step i we have the shortest paths with at most i edges.
- **Inductive Step:**
 - (First part) When $d(v)$ is updated to $d(r) + w(r, v)$, there is a path from u to r to v with weight $d(v)$.
 - (Second part) Before updating, we have by assumption the distances of all the shortest paths with at most i edges computed (if they exist).
 - Then at each possible update, we compare the shortest path to a neighbour plus an edge, with the shortest path on at most i vertices, and keep the minimum.
 - So we have a path on $i + 1$ edges that is shorter than the i edge path (and we pick the smallest of all these), or we keep the i edge path – thus the path must be the shortest path on $i + 1$ edges – there are no other possible ways to get it. \square