

31251 – Data Structures and Algorithms

Week 10, Autumn 200

Xianzhi Wang

This week, in the exciting world of algorithms:

- String Searching
- A use of hashing other than storing data
- A Probabilistic Algorithm
- Another Dynamic Programming Algorithm

String Searching

Alphabet (Σ): a finite set of symbols.

$A = (a_0, a_1, \dots, a_{n-1})$: a string over Σ , $n > 0$.

$B = (b_0, \dots, b_{m-1})$: another string over Σ , $0 < m \leq n$.

String Matching — determining whether B is a substring of A .

That is to tell whether there is some $k (\geq 0)$ such that

$$b_0 = a_k,$$

$$b_1 = a_{k+1},$$

$\dots,$

$$b_{m-1} = a_{k+m-1}.$$

► Example

A Naïve Deterministic Solution

We can take a sliding window approach:

- ① Start with B lined up with the beginning of A .
- ② Match the characters one by one.
 - If it matches, we've found it.
 - Otherwise, move B along one character and start again.
- ③ Stop when we find it, or run out of A to check against.

There's $n - m + 1$ positions that B could start at, and we do m comparisons each time, so we get $O((n - m + 1)m) \leq O(mn)$.

A Slightly Better Approach

Replace strings of length m with a function $f : \Sigma^m \rightarrow \mathbb{N}$ - then we only need to do one comparison at each step.

- We compute $\beta = f(b_0, \dots, b_{m-1})$ and then

$$\begin{array}{rcl} \alpha_0 & = & f(a_0, \dots, a_{m-1}) \\ \alpha_1 & = & f(a_1, \dots, a_m) \\ \vdots & \vdots & \vdots \\ \alpha_{n-m} & = & f(a_{n-m}, \dots, a_{n-1}) \end{array}$$

and compare only the substrings when $\alpha_j = \beta$.

The question is what function f to use:

- We could just take $f(x_0, x_1, \dots, x_{m-1}) = \sum_{i=0}^{m-1} x_i$.
 - This is pretty easy to compute [▶ Example](#).
 - However too many strings have the same value under f - too many “collisions”.
- We can pick better f and get a better result.

We can use an algorithm that employs the idea of [hashing](#).

Hashing produces a “fingerprint” for each string.

- We can pick a hash function such that any two different strings will probably produce a different hash.

Pick a large prime p and a random integer $r \in [1, p - 1]$,

We can set

$$f(x_1, x_2, \dots, x_m) = \sum_{i=\{1,2,\dots,m\}} (x_i r^{m-i}) \bmod p$$

or

$$f(x_0, x_1, \dots, x_{m-1}) = \sum_{i=\{0,1,\dots,m-1\}} (x_i r^{m-1-i}) \bmod p$$

and compute this efficiently.

How Many Collisions Do We Have Now?

How often do we have collisions, i.e.,

$$f(c_0, \dots, c_{m-1}) == f(d_0, \dots, d_{m-1}) ?$$

Not too often because the above collision implies that

$$(e_0 r^{m-1} + e_1 r^{m-2} + \dots + e_{m-2} r + e_{m-1}) \bmod p == 0$$

where $e_i = c_i - d_i$.

According to Lagrange Theorem¹, there are at most $m - 1$ “bad” values of r causing collisions, for each pair of m -tuples $(a_j, \dots, a_{n-j+1}) \neq (b_0, \dots, b_{m-1})$.

Considering all comparisons, we have at most

$$\underbrace{(m-1)}_{\text{collision\# per pair}} * \underbrace{(n-m+1)}_{\text{\# of pairs}}$$

collisions.

If $p \gg (m-1)(n-m+1)$, collision is very unlikely.

¹A polynomial of degree k has at most k roots.

- So we just choose p and r , and then compute f for all length m substrings.
- We only check if a substring (a_j, \dots, a_{j+m-1}) is equal to B when $f(\alpha_j) == f(\beta)$.
- Otherwise B is not a substring of A .

► Example

How do we compute f efficiently?

We can adapt the sliding window approach

- For overlapping substrings, we can reuse previous results -
Dynamic Programming!

$$\begin{aligned} f(a_{j+1}, \dots, a_{j+m}) &= a_{j+1}r^{m-1} + \dots + a_{j+m} \\ &= r(a_{j+1}r^{m-2} + \dots + a_{j+m-1}) + a_{j+m} \\ &= r(f(a_j, \dots, a_{j+m-1}) - a_jr^{m-1}) + a_{j+m} \end{aligned}$$

As such, we can compute all the f values for A in **linear** time!

With high probability we only have to do **m** comparisons.

This is also a probabilistic algorithm!

- It is still possible to get a lot of collisions.
- But if we start with too many, we can just guess a different r and start again.

What does that all mean practically?

- With high probability, the algorithm finds the substring in $O(n + m) = O(n)$.
 - Better than the naïve $O(nm)$ algorithm.
 - Central to this is computing the hash fingerprint quickly – reusing previous results is key.
- In the worse case, we end up with $O(nm)$.

Examples

Example

Given the alphabet $\Sigma = \{*, \&, \%\}$
and the string

$A = \& * \& \% * \% * * \& * \& * \% \% * \% * * \& \% * \underbrace{\& * * \%}_{B} \& *$

If $B = \& * * \%$ then Yes!

If $B = \% * * \%$ then No!

► Back

Using a Function to Match Strings

Example

$$\Sigma = \{*, \&, \%\} \longrightarrow \Sigma = \{0, 1, 2\}$$

$$A = \& * \& \% * \% * * \& * \& * \% \% * \% * * \& \% * \& * * \% \& *$$

$$A = 101202001010220200120100210$$

$$B = \& * * \% \longrightarrow B = 1002$$

$$\beta = 3$$

$$\alpha_j = 4, \textcolor{red}{3}, 5, 4, 2, \textcolor{red}{3}, 1, 2, 2, \textcolor{red}{3}, 5, 4, \\ 6, 4, 2, \textcolor{red}{3}, \textcolor{red}{3}, \textcolor{red}{3}, 4, \textcolor{red}{3}, 1, \textcolor{red}{3}, \textcolor{red}{3}, \textcolor{red}{3}$$

Example

$$A = 101202001010220200120100210 \quad B = 1002$$

$$n = \text{length}(A) = 27$$

$$m = \text{length}(B) = 4$$

$$n - m + 1 = 27 - 4 + 1 = 24$$

Choose $p = 9973$, $r = 5347$, Then

$$\beta = 1258$$

$$\alpha_j|_{j=0,1,\dots,24} = \{ \begin{array}{l} 6605, 8512, 6867, 3233, 5609, 2513, \\ 5347, 7792, 6603, 7793, 1979, 6330, \\ 8123, 3233, 5609, 2513, 5349, 8512, \\ 6866, 7859, 7791, \textcolor{red}{1258}, 722, 983 \end{array} \}$$